



SOFTWARE-DEFINED NETWORKING SESSION I

Introduction to Software-defined Networking
Block Course – 16-20 March 2015

David Koll

Course overview

Block Course Software-defined Networking



Introduction: basics (this week)

Advanced: more sophisticated topics
(next week)

Formalities:

- 5 ECTS for this week
- Requirements to pass:
 - 50% of points in exercises
 - Presentation of research papers (graded)
 - Written report reviewing the course (graded)
- Exercise submission:
 - All exercises together with report (some earlier)
 - Deadline: April 30th, 23:59 CET
 - Submit code as well!

IMPORTANT:
FlexNow registration
deadline:
TODAY, 23:59 CET

Course overview

Daily Schedule:

Session	Time Slot
Morning Lecture	09.15-10.45
Morning Exercise	11.15-12.45
Afternoon Lecture	14.00-15.30
Afternoon Exercise	16.00-17.00

We will deviate from this schedule at some points throughout the course

Detailed Course overview

Day	Morning Session 1	Morning Session 2	Afternoon Session 1	Afternoon Session 2
Mon	Introducing SDN I	Lecture Exercises	Introduction to Python	Python Exercises
Tue	Introducing SDN II	Lecture Exercises	Introduction to Python (cont.)	Python Exercises
Wed	Current Research in SDN	Paper Selection and Reading (Teams)	Hands-On SDN I	SDN Exercises
Thu	Hands-On SDN II	SDN Exercises	Hands-On SDN III	Presentation Preparation
Fri	Hands-On SDN IV	SDN Exercises	Wrap-Up & Free Slot	Presentations

Detailed Course overview

Day	Morning Session 1	Morning Session 2	Afternoon Session 1	Afternoon Session 2
Mon	Introducing SDN I	Lecture Exercises	Introduction to Python	Python Exercises
Tue	Introducing SDN II	Lecture Exercises	Introduction to Python (cont.)	Python Exercises
Wed	Current Research in SDN	Paper Selection and Reading (Teams)	Hands-On SDN I	SDN Exercises
Thu	Hands-On SDN II	SDN Exercises	Hands-On SDN III	Presentation Preparation
Fri	Hands-On SDN IV	SDN Exercises	Wrap-Up & Free Slot	Presentations

Lectures: Basic concepts you need to know about SDN and programming for SDN

Detailed Course overview

Day	Morning Session 1	Morning Session 2	Afternoon Session 1	Afternoon Session 2
Mon	Introducing SDN I	Lecture Exercises	Introduction to Python	Python Exercises
Tue	Introducing SDN II	Lecture Exercises	Introduction to Python (cont.)	Python Exercises
Wed	Current Research in SDN	Paper Selection and Reading (Teams)	Hands-On SDN I	SDN Exercises
Thu	Hands-On SDN II	SDN Exercises	Hands-On SDN III	Presentation Preparation
Fri	Hands-On SDN IV	SDN Exercises	Wrap-Up & Free Slot	Presentations

Exercises: Review and/or implement the concepts taught in the lectures

Detailed Course overview

Day	Morning Session 1	Morning Session 2	Afternoon Session 1	Afternoon Session 2
Mon	Introducing SDN I	Lecture Exercises	Introduction to Python	Python Exercises
Tue	Introducing SDN II	Lecture Exercises	Introduction to Python (cont.)	Python Exercises
Wed	Current Research in SDN	Paper Selection and Reading (Teams)	Hands-On SDN I	SDN Exercises
Thu	Hands-On SDN II	SDN Exercises	Hands-On SDN III	Presentation Preparation
Fri	Hands-On SDN IV	SDN Exercises	Wrap-Up & Free Slot	Presentations

Presentations: train your presentation skills based on hot topics in SDN (20+5)

Final report

Grading:

- 100% final report
- LaTeX template will be provided on course wiki page

[https://wiki.net.informatik.uni-goettingen.de/wiki/Introduction_to_Software-defined_Networking_\(Winter_2014/2015\)](https://wiki.net.informatik.uni-goettingen.de/wiki/Introduction_to_Software-defined_Networking_(Winter_2014/2015))

Final report

Report has to contain:

- Reviews of the research papers you read (review form will be available)
- Your code from the exercises
- Submit in an archive to:

`dkoll~at~cs.uni-goettingen.de`

WRITING DOS AND DON'TS

- **DO:** use your own words!
- **DO:** incorporate your own ideas
- **DO:** cite properly
- **DO:** proofread it!
- **DON'T:** plagiarize (please do not copy from the papers you read)

[50] Kurt Thomas, Damon McCoy, Chris Grier, Alek Kolcz, and Vern Paxson. Trafficking Fraudulent Accounts: The Role of the Underground Market in Twitter Spam and Abuse. In *Proceedings of the 22nd USENIX Security Symposium (USENIX Security 2013)*, pages 195–210. USENIX, 2013.

Introduction to Software-defined Networks – Session I

**Partly based on slides of Nick McKeown, Scott Shenker, Nick
Feamster, and Jennifer Rexford**

Why this course?

„Software-Defined Networks – **the counter model of the internet**“
– *heise.de*

“**November 2014: Cisco** declares “game over” for SDN competitors [...], prompting reaction from two industry groups that the game has just begun; **Alcatel-Lucent** and **Juniper** also virtualize their routers [...]; **AT&T** and others unveil [...] an alternative [...].”
– *networkworld.com*

“Many solution providers believe 2015 is the year that **SDN will truly begin to reshape the networking landscape**”
– *crn.com*

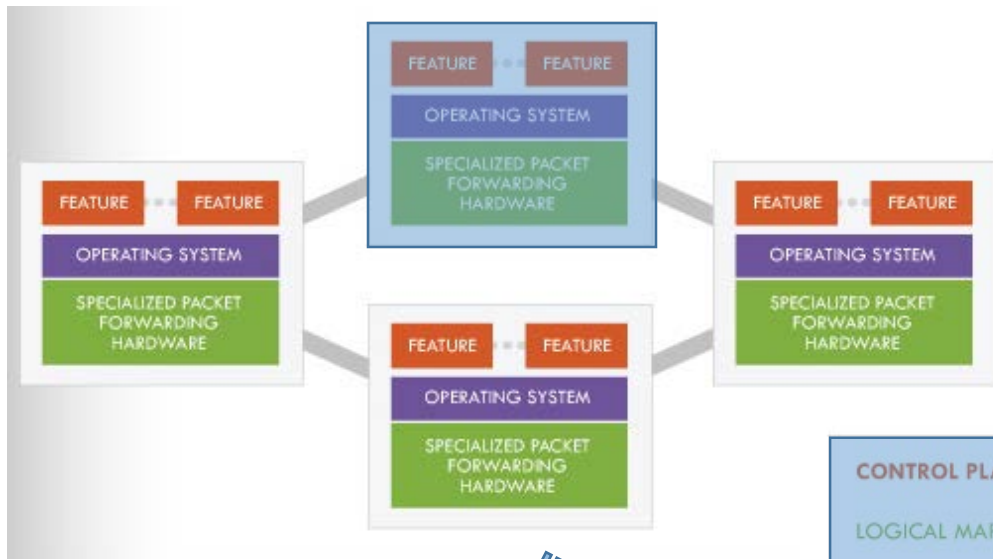
What is Software-defined Networking?

“The physical separation of the network control plane from the forwarding plane, and where a control plane controls several devices.”

– **The Open Networking Foundation***

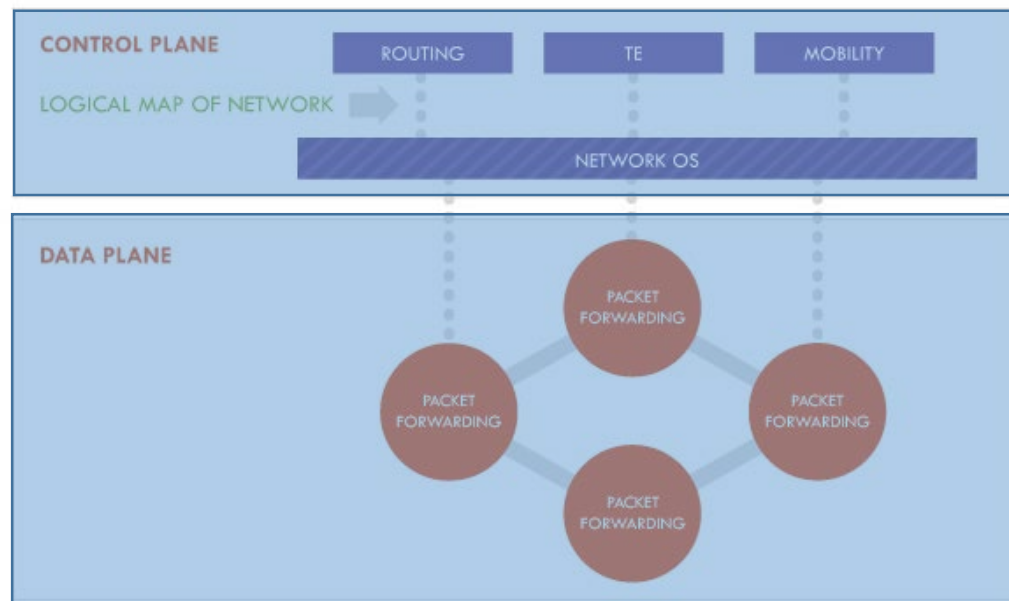

*** Google, Facebook, Microsoft, Deutsche Telekom, Verizon, Yahoo, Cisco, Citrix, Dell, Ericsson, HP, IBM, Juniper Networks, NEC, Netgear, VMWare, ...
...and various institutions from academia (e.g., Stanford, Berkeley)**

SDN in a Nutshell



“The physical separation of **the network control plane** from the **forwarding plane**, and where a **control plane controls several devices.**”

– The ONF



Taken from: <http://www.opennetsummit.org/archives/apr12/site/why.html>

The History behind the Hype

Going to talk about...

What are the origins of SDN?

Why do we need SDN?

Where are we now?

... before we dive into the
technical details of SDN


The History behind the Hype

The concepts behind SDN are not really new!

Scott Shenker: „*[SDN is] not a revolutionary technology, [it is] just a way of organizing network functionality.*“[1]


[1] S. Shenker in his talk „A Gentle Introduction to Software-defined Networks“

The History behind the Hype



N. Fearnster et al.: „The Road to SDN – An intellectual history of programmable networks“ ACM SIGCOMM Computer Communication Review 44.2 (2014): 87-98.

The History behind the Hype



N. Fearnster et al.: „The Road to SDN – An intellectual history of programmable networks“ ACM SIGCOMM Computer Communication Review 44.2 (2014): 87-98.

A brief history of programmable networks: Active Networks

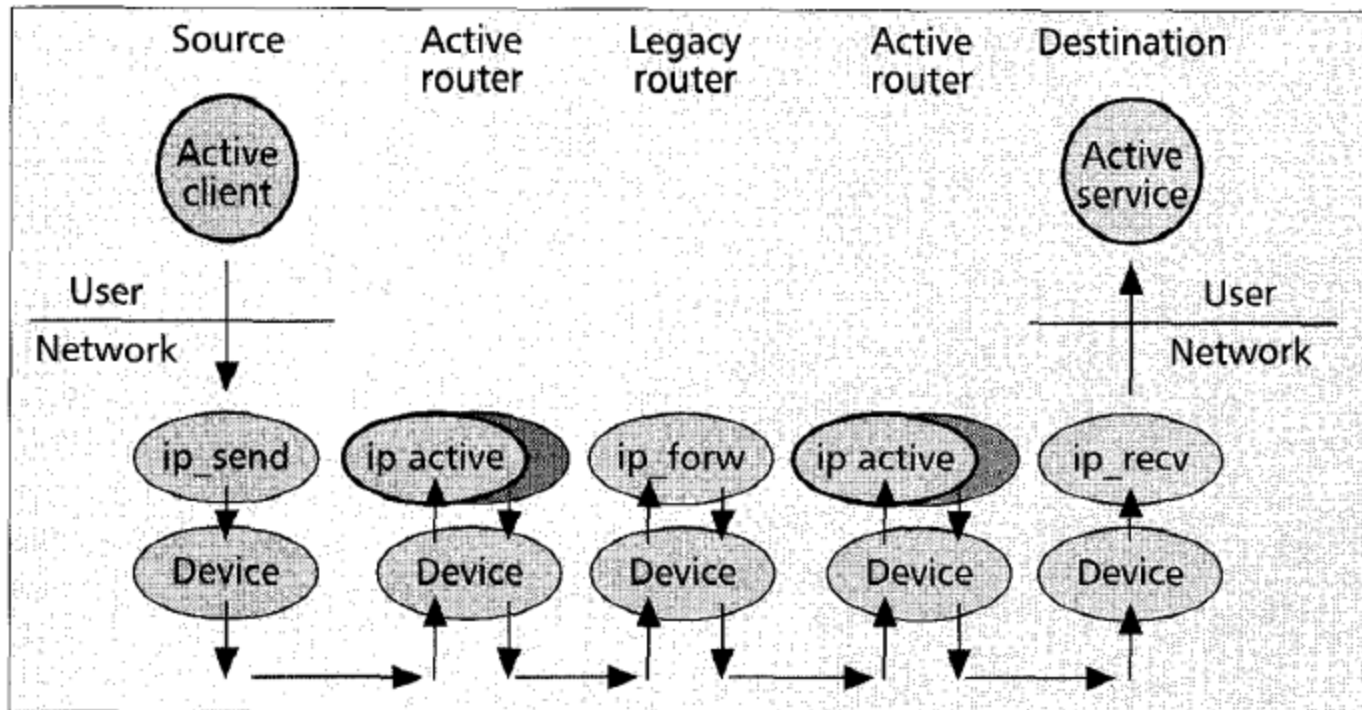


Active Networks?

- End of 1990s: network ossification (idea->deployment: 10 years!)
- Goal: opening up network control
- Envisioned method: make network devices programmable via an API
- API could be accessed via two models:
 - **Capsule model:** code included in data packets transmitted in-band [1]
 - **Programmable router/switch model:** code transmitted out-of-band [2]

[1] Wetherall, et al.: "ANTS: a toolkit for building and dynamically deploying network protocols." In Proceedings of IEEE OpenArch 1998.
[2] Bhattacharjee, S., Calvert, K.L. et al.: "An architecture for active networks". In Proceedings of High-Performance Networking 1997.

Active Networks



■ **Figure 1.** *Application-specific processing within the nodes of an active network.*

Co-existence of legacy routers with active routers

[1] Tennenhouse, et al.: "A survey of active network research." *IEEE Communications Magazine*, 35.1 (1997): 80-86.



Why did active networks fail?

- Timing was off
 - End of 1990s: no data-centers/clouds yet
 - Hardware was expensive (compared to 2015)
- Conceptual mistakes:
 - Programmable by end-users (security?)
 - Limited interoperability



The Legacy of Active Networks

- Intellectual contributions of Active Networks:
 - Programmable network functions
 - Network virtualizations (de-multiplexing of packets according to their header)

The concepts behind SDN are not really new!
(we see both contributions in today's SDN)

- [1] Wetherall, et al.: "ANTS: a toolkit for building and dynamically deploying network protocols." In Proceedings of IEEE OpenArch 1998.
[2] Bhattacharjee, Calvert, et al.: "An architecture for active networks". In Proceedings of High-Performance Networking 1997.

A brief history of programmable networks: Control and data plane separation



Control and Data Plane Separation

- Early 2000s: increasing traffic volumes, network sizes
 - need for traffic engineering
- But: conventional routers/switches: tight integration of data and control planes
 - Problem: Hard to debug and control router behaviour
- Goal: Traffic control and configuration should be easier
- Envisioned method: decouple control and data plane



Control and Data Plane Separation

- Mainly two innovations:
 - Open interface between the control and data plane (e.g., ForCES [1])
 - *Logically* centralized control of the network (e.g., RCP [2])
- Compared to active networks:
 - **Targeted at network administrators** rather than end-users
 - **Programmability in control plane** rather than in data plane
 - **Network wide control** rather than device-level configuration


[1] Yang, et al. "Forwarding and control element separation (ForCES) framework." RFC 3746, April, 2004.

[2] Caesar, et al. "Design and implementation of a routing control platform." Proceedings of Usenix NSDI, 2005.

RCP - Separating *Interdomain* Routing [1]

- Compute interdomain routes for the routers
 - Input: BGP-learned routes from neighboring ASes
 - Output: forwarding-table entries for each router
- Backwards compatibility with legacy routers
 - RCP speaks to routers using BGP protocol
- Routers still run intradomain routing protocol
 - So the routers can reach the RCP
 - To reduce overhead on the RCP

**Autonomous
System (AS)**




[1] Caesar, et al. "Design and implementation of a routing control platform." Proceedings of Usenix NSDI, 2005.

Incremental Deployability


- Backwards compatibility
 - Work with existing routers and protocols
- Incentive compatibility
 - Offer significant benefits, even to the first adopters
 - E.g., reducing overhead at routers

~~AS 1, AS 2, AS 3 are not connected to each other. AS 1 is connected to AS 2 and AS 2 is connected to AS 3. AS 1 is connected to AS 2 and AS 2 is connected to AS 3. AS 1 is connected to AS 2 and AS 2 is connected to AS 3.~~




Example: Maintenance Dry-Out

- Planned maintenance on an edge router
 - Drain traffic off of an edge router
 - *Before* bringing it down for maintenance



Example: Egress Selection

- Customer-controlled egress selection
 - Multiple ways to reach the same destination
 - Giving customers control over the decision



RCP – The big “BUT”

- RCP still uses BGP, a single routing protocol
 - This is not what we need
- However, we can learn from it!



The Legacy of the Separation

- Recall the two innovations:
 - Open interface between the control and data plane (e.g., ForCES [1])
 - *Logically* centralized control of the network (e.g., RCP [2])

The concepts behind SDN are not really new!
(we see both contributions in today's SDN)


[1] Yang, et al. "Forwarding and control element separation (ForCES) framework." RFC 3746, April, 2004.

[2] Caesar, et al. "Design and implementation of a routing control platform." Proceedings of Usenix NSDI, 2005.

A brief history of programmable networks: Network virtualization

Network Virtualization?

- Abstraction of a network that is decoupled from the underlying physical network (e.g., VLANs)



Server virtualization

- Run multiple virtual servers on a physical server
- Each VM has illusion it is running as a physical server

Network virtualization

- Run multiple virtual networks on a physical network
- Each virtual network has illusion it is running as a physical network

Microsoft Technet.: <https://gallery.technet.microsoft.com/scriptcenter/Simple-Hyper-V-Network-d3efb3b8>



Network Virtualization

First steps:

- Overlay networks as virtual networks on top of legacy technology
 - Own control protocol, encapsulation over legacy network (tunneling)
 - MBone [1] (for multicast), 6Bone [2] (for IPv6)

**In contrast to active networks, overlay networks
do not require any support from network equipment**

Later:

Virtual networks inside the underlying network (e.g., VINI [3])

[1] Almeroth et al, "Multicast group behavior in the Internet's multicast backbone (MBone)." *IEEE Communications Magazine*, IEEE35.6

[2] Fink et al, *6bone (IPv6 testing address allocation) phaseout*. RFC 3701, March, 2004.

[3] Bavier, et al. "In VINI veritas: realistic and controlled network experimentation." ACM CCR. Vol. 36. No. 4. ACM, 2006.

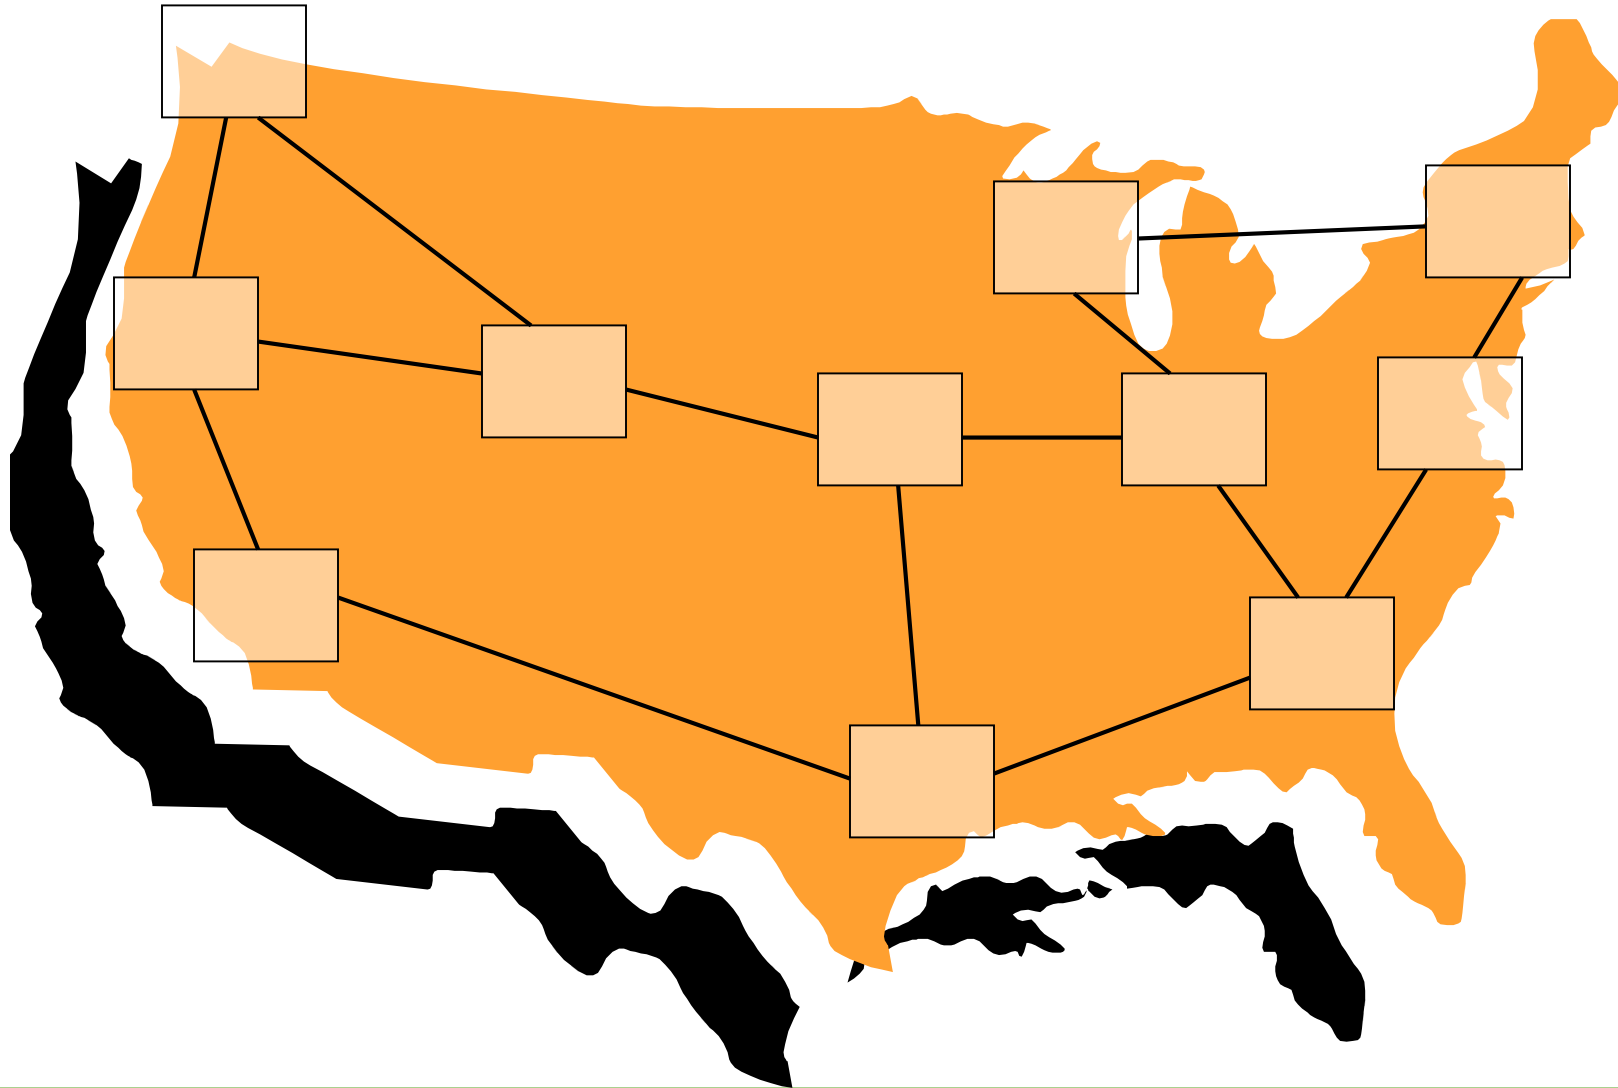


How to Validate an Idea?



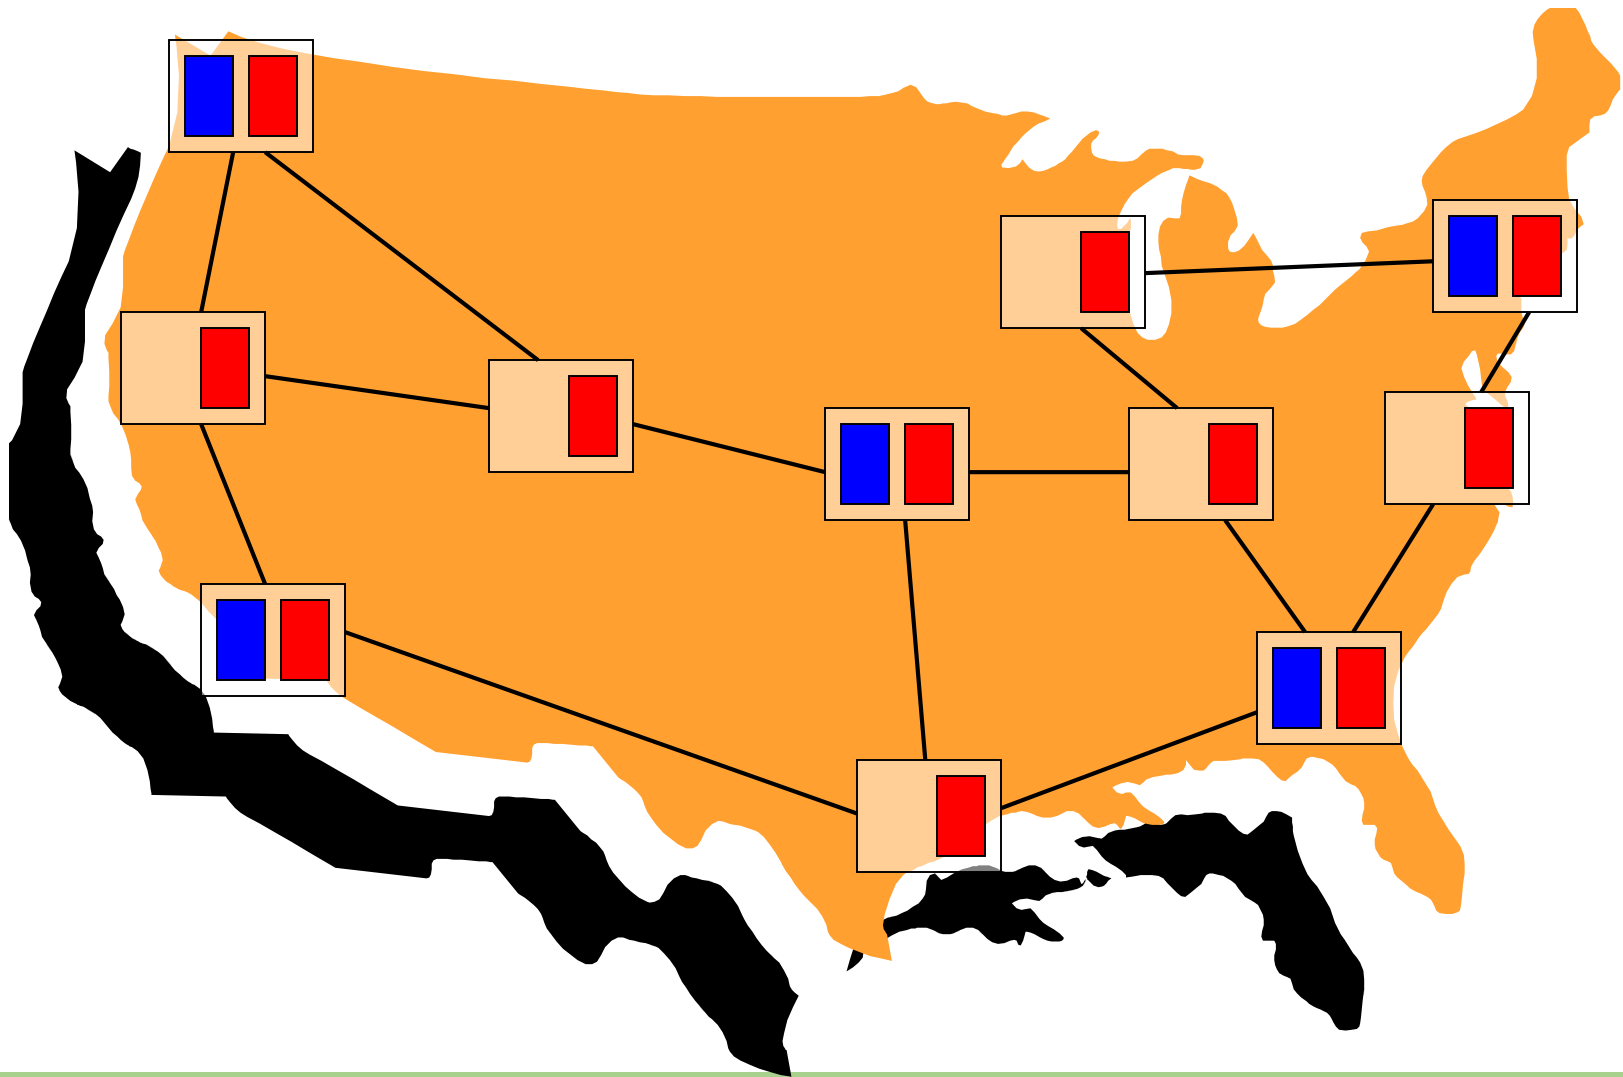
- Fixed, shared among many experiments
- Runs real routing software
- Exposes realistic network conditions
- Gives control over network events
- Carries traffic on behalf of real users

Fixed Infrastructure

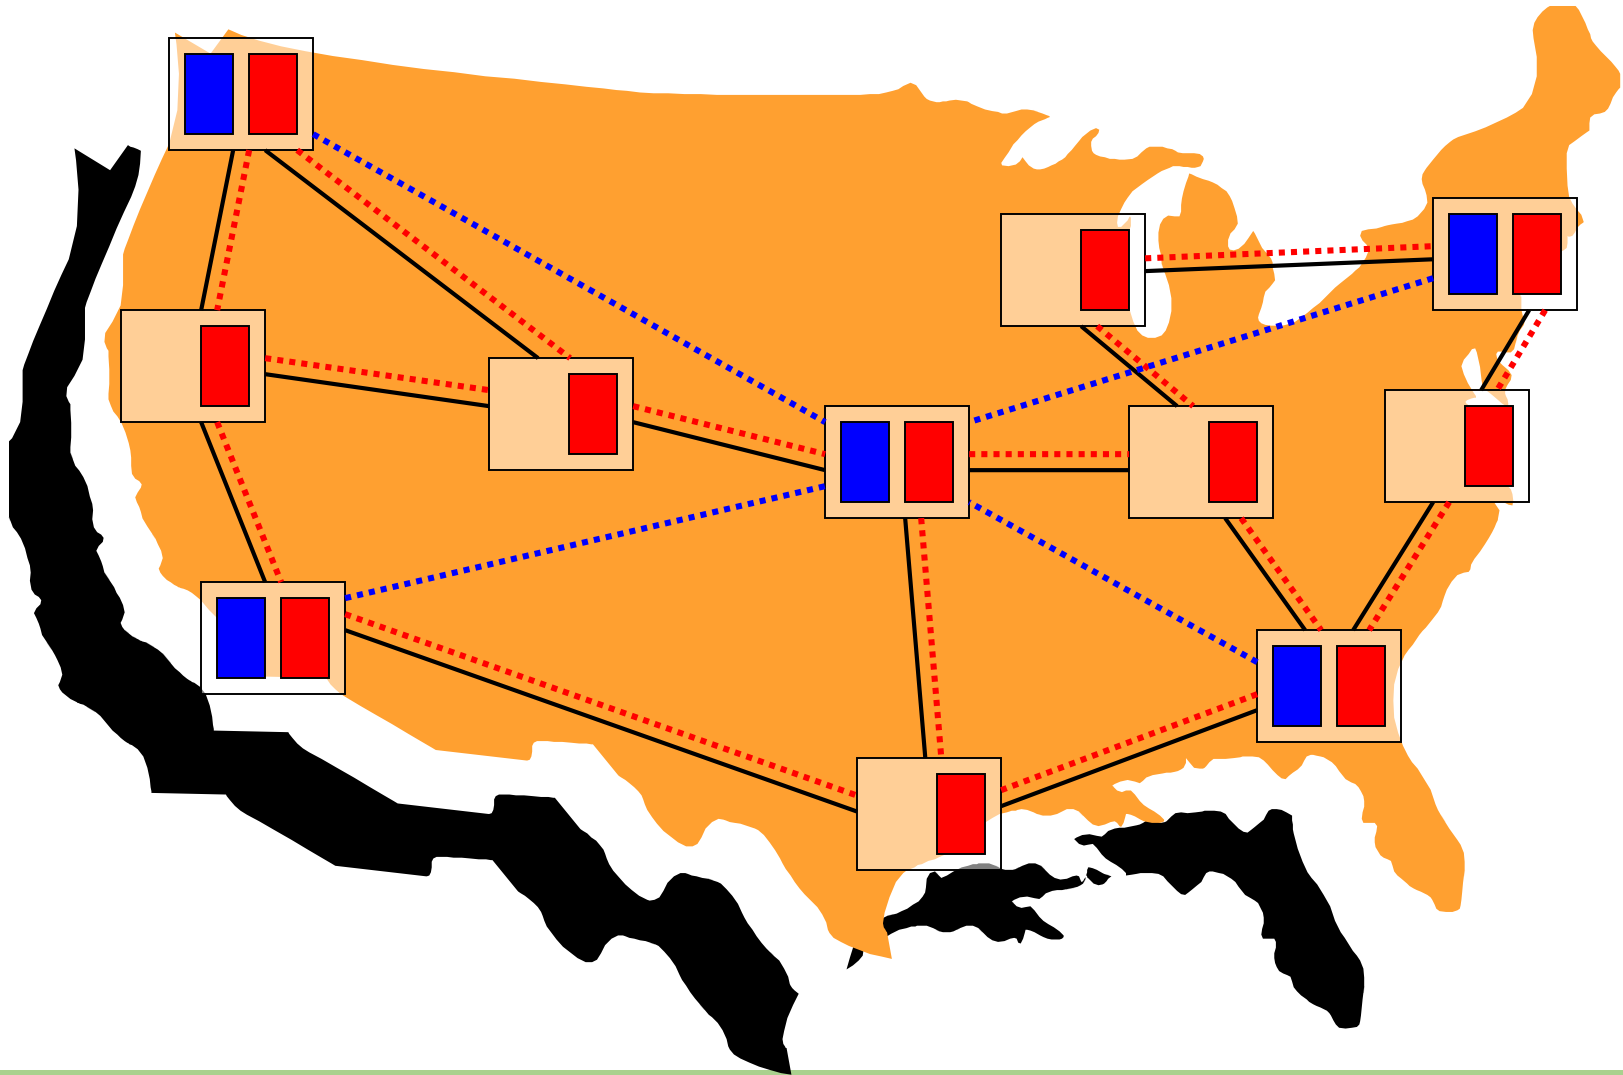


[3] Bavier, "VINI and its future directions"
http://www.fp7-federica.eu/pres_eventi/VINI_TNC2008.ppt

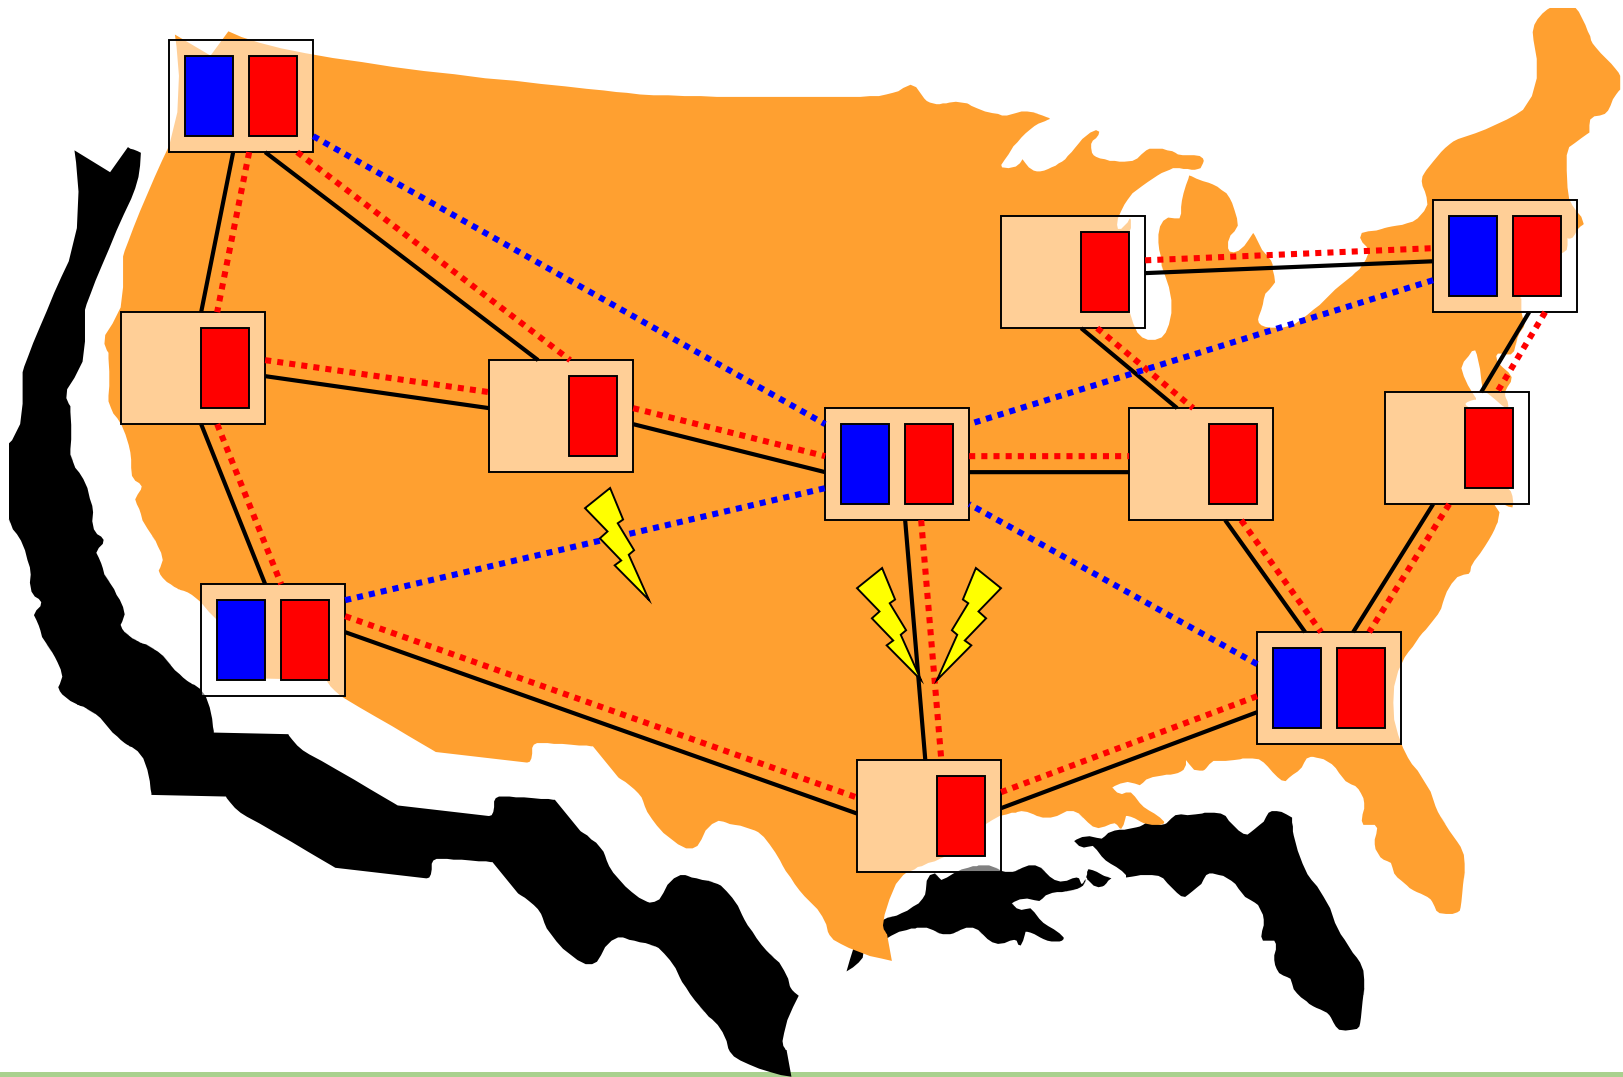
Shared Infrastructure



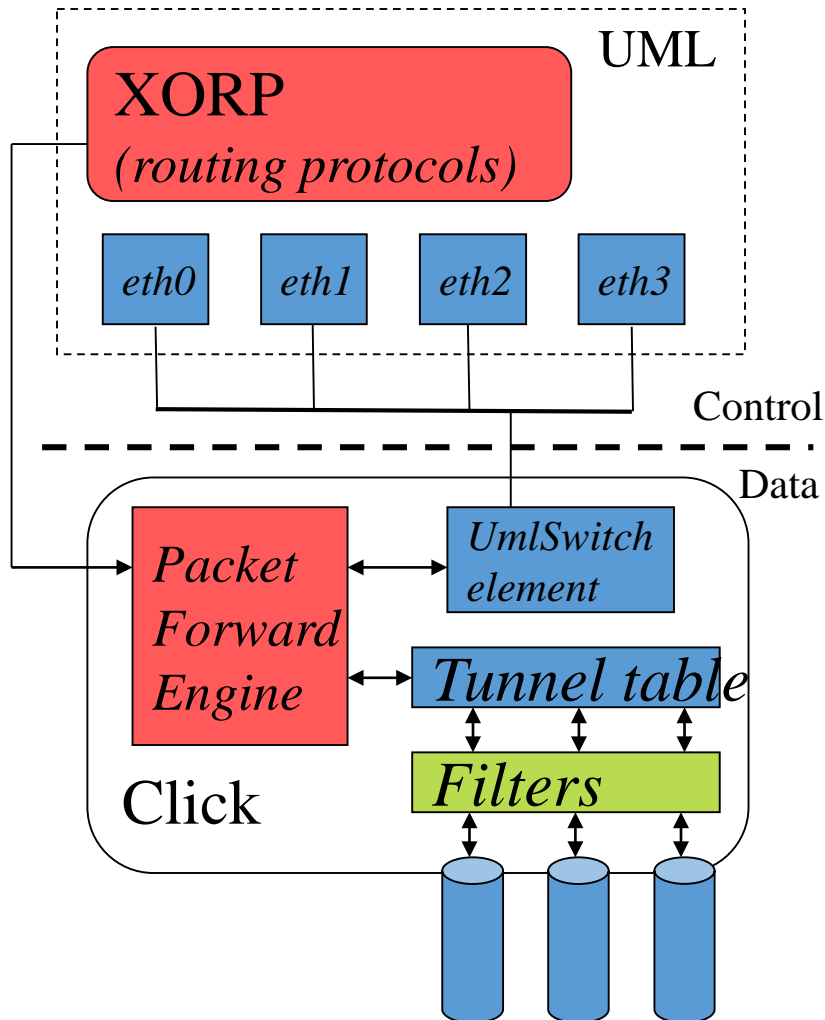
Flexible Topology



Network Events



VINI: Control/Data Plane Separation



- Performance
 - Avoid UML overhead
 - Move to kernel, FPGA
- Interfaces \Rightarrow tunnels
 - Click UDP tunnels correspond to UML network interfaces
- Filters
 - “Fail a link” by blocking packets at tunnel



The Legacy of Network Virtualization

- Three main ideas
 - Separate service from infrastructure
 - Have multiple controllers (virtual networks) for the same switch
 - Logical network topologies

The concepts behind SDN are not really new!
(we see these contributions in today's SDN)

Control and Data Plane Separation – What *is* SDN actually?

SDN: Control and Data Plane Separation

Control Plane

logic for controlling the forwarding elements
routing protocols (e.g., BGP, OSPF), middlebox configuration, etc.

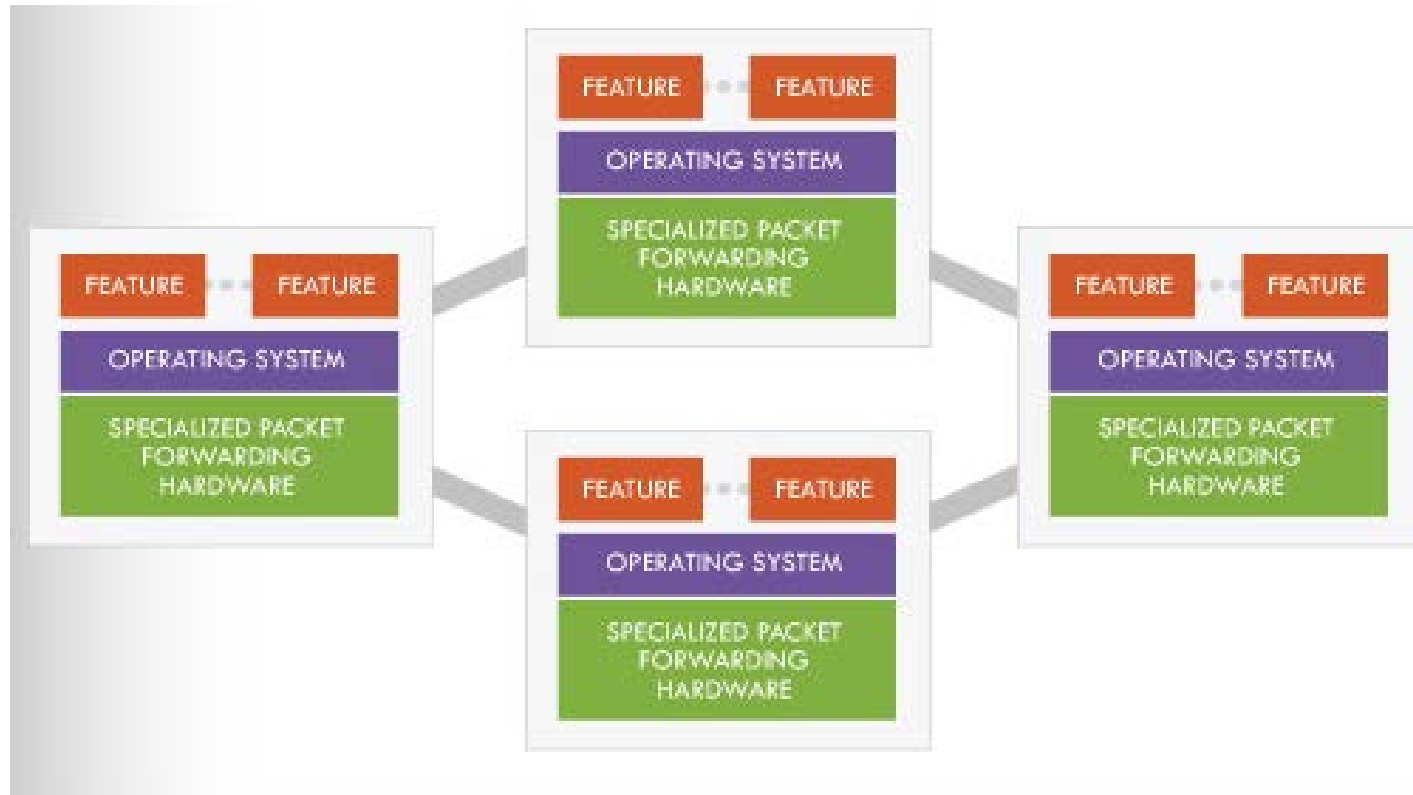
Data Plane

forward data based on rules set by the control logic
IP forwarding, layer 2 switching, etc.

Today, routers implement both

Why separate?

Currently, routers implement *both*:



What do we gain from separating control and data plane?

Key to Internet Success: Layers

Applications

...built on...

Reliable (or unreliable) transport

...built on...


Best-effort global packet delivery

...built on...

Best-effort local packet delivery

...built on...

Physical transfer of bits




Why Is Layering So Important?

- It provides **abstraction**: decomposed delivery into fundamental components
- Independent but compatible innovation at each layer
- A practical success (it still works!)

The Problem in Computer Networks

- Layers only deal with the **data plane**
- We have no powerful ***control plane*** *abstractions!*

Control Plane Abstractions




The Problem in Computer Networks

- Many different control plane mechanisms
- **Designed from scratch for specific goal**
- Variety of implementations
 - **Globally distributed:** routing algorithms
 - **Manual/scripted configuration:** ACLs, VLANs
 - **Centralized computation:** Traffic engineering
- **Network control plane is a complicated mess!**

The Problem in Computer Networks

- Complexity has increased to “unmanageable” levels
- Consider datacenters:
 - 100,000s machines, 10,000s switches
 - 1000s of customers
 - Each with their own logical networks: ACLs, VLANs, etc
- Way beyond what we can handle
 - Leads to brittle, ossified configurations
 - Inefficient as well

Example: Datacenter Networks



Example: Datacenter Networks

- Complexity has increased to “unmanageable” levels
 - | **20k server cluster \approx 16k internal links**
 - | This means upto **1024** distinct links between a pair of hosts
 - | How do you troubleshoot this (for packetloss, etc)?
 - | # of links to test = $1024/2 = 512$
 - | 30 seconds/test
 - | **256 man-minutes for most-basic troubleshooting!**

The Problem is not only Complexity

- Closed equipment
 - Software bundled with hardware
 - Vendor-specific interfaces
- Over specified
 - Slow protocol standardization
- Few people can innovate
 - Equipment vendors write the code
 - Long delays to introduce new features




Enter SDN

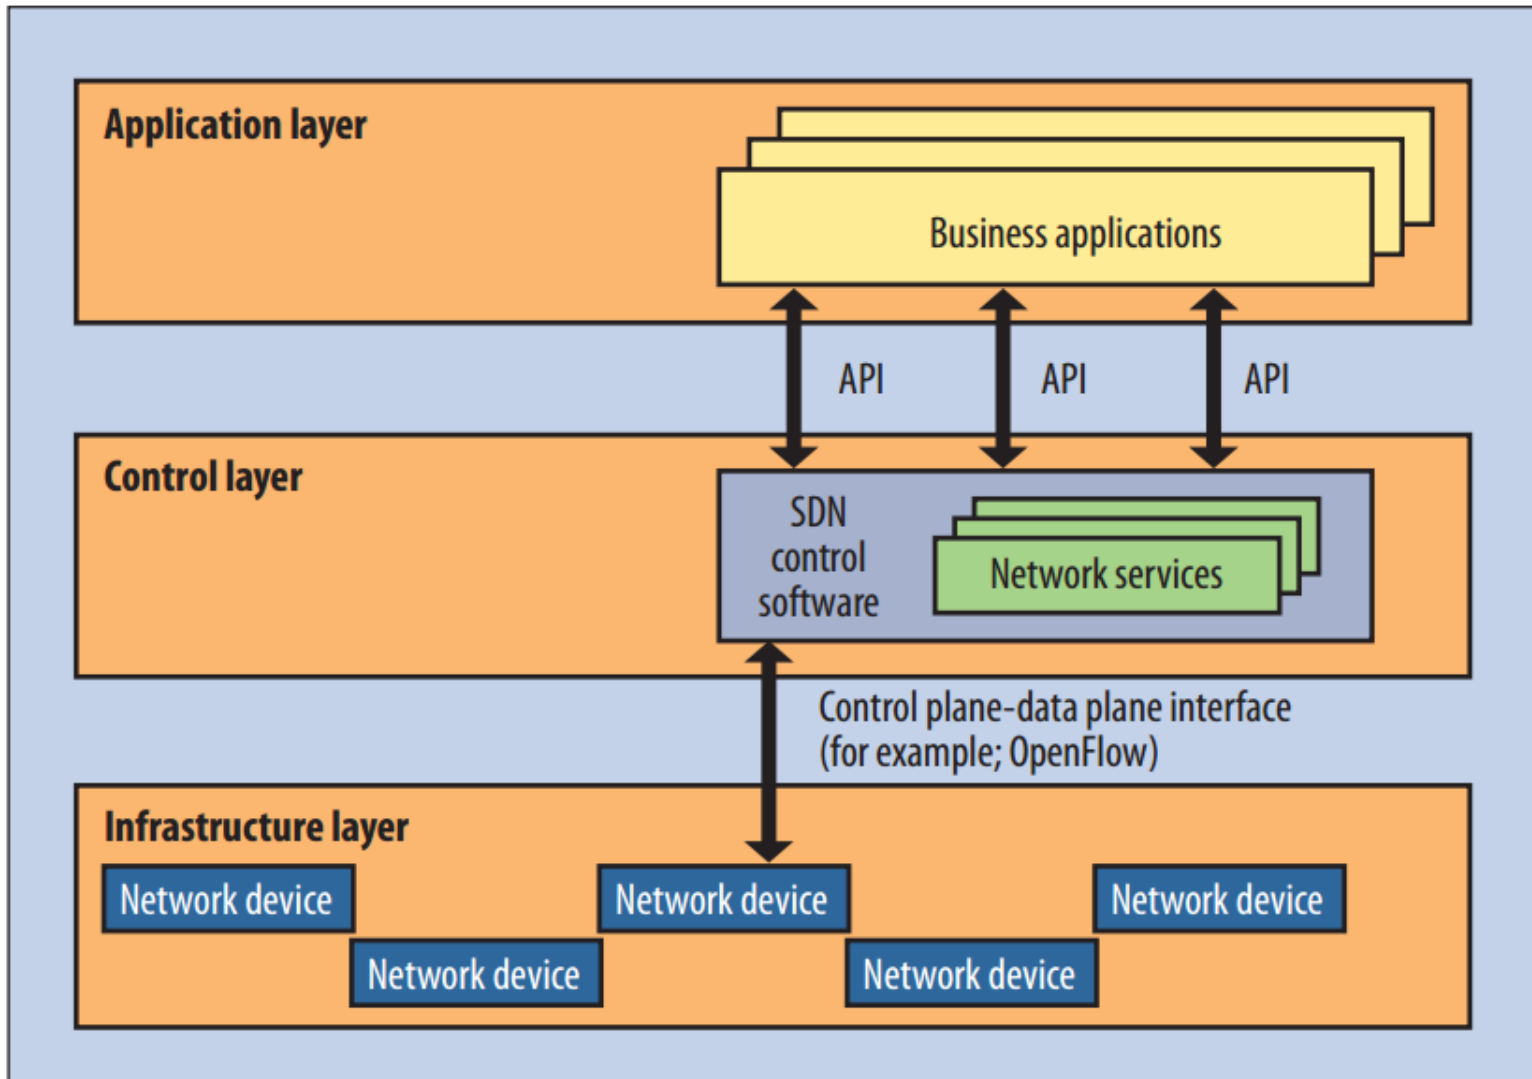
- Today, routers implement both planes
 - They forward packets
 - And run the control plane software
- SDN networks
 - Data plane implemented by switches
 - Switches act on local forwarding state
 - Control plane implemented by controllers
 - All forwarding state computed by SDN platform
 - Open protocols!
- **A technical change with broad implications**

Enter SDN

*e.g. routing,
access control*



Another View



<http://www.networkcomputing.com/networking/searching-for-an-sdn-definition-what-is-software-defined-networking/>

Anology

- You are lost in a city and are trying to reach a destination
- Today's networks: ask other people you meet to obtain information (routing protocols)
- SDN: pull out your cellphone and start Google maps – it will calculate the route for you

Changes

- Less vendor lock-in
 - Can buy HW/SW from different vendors
- Changes are easier
 - Can test components separately
 - *HW has to forward*
 - *Can simulate controller*
 - *Can do verification on logical policy*
 - Can change topology and policy independently
 - Greater rate of innovation

Challenges of the Separation

- Talked a lot about the opportunities
- What about the challenges?

Practical Challenges

- Scalability
 - Control elements responsible for many routers
- Response time
 - Delays between control elements and routers
- Reliability
 - Surviving failures of control elements and routers
- Consistency
 - Ensuring multiple control elements behave consistently
- Security
 - Network vulnerable to attacks on control elements
- Interoperability
 - Legacy routers and neighboring domains

Example - Scalability

- Take routing: the controller has to make routing decisions for a lot of routers
 - Potentially 1000s
- Also has to store these routes
 - a lot of routing tables
- Single controller node for this task?
 - Compare with current standard OSPF: distributed

Current Status of SDN

- SDN widely accepted as “**future of networking**”
 - ~1000 engineers at latest Open Networking Summit
 - *Much more acceptance in industry than in academia*
- Insane level of SDN hype, and still:
 - SDN doesn't work miracles, merely makes things easier

Current Status of SDN

- Most innovations in southbound interface, controllers, northbound interface, and applications
 - OpenFlow (as ONE example of the sb interface)
 - NOX, POX, ONOS, etc.
 - Pyretic, Frenetic, etc.
- But: also changes in network devices
 - Most global players offer SDN switches now

Up Next

